

## AlphaGo: Mastering the game of Go with deep neural networks and tree search

Solving the game of Go has been a long standing challenge in the realm of artificial intelligence. As a two-person game of perfect information, it can in principle be solved exactly. However, due to the game's large branching factor ( $b \approx 250$ ) and depth ( $d \approx 150$ ), searching through the entire space of possible moves ( $b^d \approx 10^{359}$ ) is unfeasible. Instead, one could perform a depth-limited search by approximating the optimal value function  $v^*(s)$  of any given state of the game  $s$  by some approximate value function that truncates the search tree at the given state  $s$ , i.e.  $v^*(s) \approx v(s)$ . Further, one could prune the search tree and limit the breadth of the search by sampling random actions  $a$  given a state of the game  $s$  via some policy (conditional) probability distribution  $p(a|s)$ . This is the approach used in AlphaGo.

The architecture of AlphaGo consists of several neural networks, each designed to perform a specific task. In the first stage, a supervised machine learning approach is adopted, and a convolutional neural network with weights  $\sigma$  is trained on expert human moves to produce a probability distribution  $p_\sigma(a|s)$  of likely actions given a state of the game. For the training of this neural net, random state-action pairs  $(s, a)$  are sampled as inputs, and the minimum of the objective function is sought using stochastic gradient ascent, which maximizes the log-likelihood probability distribution as a function of the weights. This 13-layer policy neural network was able to predict expert human moves with an accuracy of up to 57.0% on a test set. Meanwhile, a smaller and less accurate, but faster, rollout policy network  $p_\pi(a|s)$  was similarly trained for the purposes of quickly generating state-action pairs  $(s, a)$  that could be fed into the convolutional policy network for its training.

The next stage in the pipeline consists of a reinforcement learning policy network  $p_p(s|a)$ , with an identical structure to the supervised learning policy network  $p_\sigma(a|s)$ , and with the weights initialized to the learned weights in the previous stage of the pipeline,  $p=\sigma$ . Games are then played between the current policy network and some randomly selected previous iteration of the same network, using a reward function  $r(s)$ , which is set to zero for all time steps before the terminal time step, i.e.  $r(s_t) = 0$  for all  $t < T$ , and non-zero only at the terminal time step  $T$ . The outcome at a given time-step  $t$  is given in terms of the terminal reward via  $z_t = \pm r(s_T)$ , and equal to +1 if the current player at time step  $t$  wins at the terminal time step  $T$ , and -1 if that player loses. Using this prescription, the weights of the reinforcement learning network are then updated at each time step  $t$  using stochastic gradient ascent to maximize the log-likelihood probability distribution times the reward function,  $\log p_p(s|a) * z_t$ . Upon sampling moves  $a$  from the learned probability distribution  $p_p(s|a)$ , this network was found to win 85% games against Pachi, the strongest open-source Go program at the time of writing of the paper.

The last stage in the training pipeline consists of learning an estimated reward or value function  $v^p(s)$  that predicts  $z_t$  (defined above) given the state of the game at any given time step  $t$ ,  $v^p(s) = E[z_t | s_t = s]$  using the previously learned policy to generate actions  $a_{t \dots T} \sim p$ . We can then approximate the value function by a value network  $v_\theta(s)$ , with an architecture similar to that of

policy network, and with weights  $\theta$  learned to minimize the mean square error between  $z$  and  $v_{\theta}(s)$ .

AlphaGo then combines all the above in a Monte Carlo Tree Search (MCTS) which involves selecting an action at each time step that maximizes a linear combination of a back-propagated action value, using the value network  $v_{\theta}(s)$  and the outcome at a leaf node  $z_L$  (see definition of  $z_t$  above), and another term that encourages more probable actions using the learned prior distribution  $P(s, a) = p_{\sigma}(a|s)$ , but penalizes the number of visit counts of each  $(s, a)$  pair to encourage exploration of the search tree. Playing the game of Go in this manner, AlphaGo became the first computer program in Oct 2015 to beat a human Go master in a 5-0 win.